

LegalSaathi

Product Requirements Document (PRD)

AI-Powered Legal Aid Agent for Rural India

HiDevs × Mastra Hackathon 2026 | Open Innovation Challenge

Mastra · Qdrant · Enkrypt AI · TypeScript

Document Type	Product Requirements Document (PRD)
Product	LegalSaathi — AI Legal Aid Agent
Hackathon	HiDevs × Mastra Hackathon 2026
Challenge	Open Innovation — Legal & Compliance Domain
Team Name	NovaX
Team Memers	Halona Ann Siju & Mannitha J

1. Executive Summary

LegalSaathi is a production-grade AI agent designed to provide accessible, accurate, and safe legal guidance to rural Indians in their native language. Built on Mastra for agent orchestration, Qdrant for semantic legal document retrieval, and Enkrypt AI for hallucination detection and safety guardrails, LegalSaathi enables rural citizens to understand their legal rights, check government scheme eligibility, and generate legal documents — all through a conversational interface available on Web, WhatsApp, and voice in 10+ Indian languages.

The product addresses a silent crisis: over 330 million rural Indians have no affordable access to legal help. LegalSaathi bridges this gap by putting a full legal team in every rural citizen's pocket — in their language, at zero cost, available 24/7.

2. Problem Statement

2.1 The Crisis

India's rural population faces a systematic denial of legal access. The consequences are severe and widespread:

Affected Population	330M+ rural Indians with no access to affordable legal help
Cost Barrier	A single legal consultation costs Rs. 500–2,000 — unaffordable for daily wage earners earning Rs. 300–500/day
Language Barrier	Legal information exists only in English or complex Hindi legalese — inaccessible to 70%+ of rural citizens
Awareness Gap	Citizens unaware of rights under PM-KISAN, MGNREGA, RTI Act, land acquisition laws
Consequence	Land loss, wrongful imprisonment, denial of scheme benefits, inability to file RTI applications

2.2 Root Causes

- No vernacular legal content — all authoritative legal sources are in English
- Geographic barrier — lawyers are concentrated in urban district headquarters
- Economic barrier — legal fees exceed daily income for most rural households
- Digital literacy gap — existing legal tech platforms require English proficiency
- Trust gap — rural citizens distrust formal legal systems due to past exploitation

2.3 Why Existing Solutions Fail

Current alternatives — Google Search, legal helplines, government portals — fail because they require English literacy, do not understand natural language queries in regional languages, provide generic information rather than personalized guidance, and lack safety mechanisms to prevent dangerous legal misinformation.

3. Goals and Success Metrics

3.1 Primary Goal

Provide accurate, safe, cited legal guidance to rural Indians in 10+ Indian languages through an AI agent that genuinely understands their situation and gives them actionable next steps.

3.2 Success Metrics

Metric	Target	Measurement Method
Query Resolution Rate	>85%	% of queries answered without DLSA redirect
Hallucination Rate	<5%	Enkrypt AI confidence score logs
Response Time	<3 seconds	End-to-end API latency monitoring
Languages Supported	10+	IndicTrans2 language coverage
RTI Draft Accuracy	>90%	Legal expert spot-check review
User Satisfaction	>4.2/5	In-app feedback rating
Monthly Active Users	10,000+ (post-launch)	Platform analytics

4. Target Users

4.1 User Segments

Segment	Profile	Primary Need
Rural Farmers	Age 30–60, low digital literacy, Hindi/regional language only	Land rights, scheme eligibility, acquisition compensation
Daily Wage Workers	Age 20–45, mobile-first, WhatsApp users	Labour law rights, MGNREGA wage disputes, wrongful termination
First-Gen Youth	Age 18–30, basic smartphone literacy, bilingual	Navigating legal systems for family, RTI filing, FIR guidance
Women in Rural Areas	Age 20–50, limited mobility, voice-first preferred	Domestic violence rights, property inheritance, scheme access
Small Business Owners	Age 25–50, semi-urban, Hindi/regional	GST queries, licensing, consumer protection rights

4.2 Geographic Scope

- All 28 Indian states and 8 Union Territories

- Primary focus: UP, Bihar, Rajasthan, MP, Maharashtra, Tamil Nadu, Andhra Pradesh, West Bengal
- Languages: Hindi, Tamil, Telugu, Bengali, Marathi, Kannada, Gujarati, Malayalam, Punjabi, Odia

5. User Stories

- **US-01: Government Scheme Eligibility**

As a farmer in rural Uttar Pradesh, I want to ask in Hindi whether I am eligible for PM-KISAN, so that I can receive the Rs. 6,000 annual benefit I may be missing out on.

Acceptance Criteria: Agent correctly identifies eligibility criteria, cross-references user's stated land holding and Aadhaar status, provides step-by-step enrollment guidance in Hindi.

- **US-02: Land Rights Query**

As a tribal farmer in Jharkhand whose land was acquired by the government, I want to understand my legal rights and compensation entitlements in my language, so that I can take the correct legal action before the deadline passes.

Acceptance Criteria: Agent retrieves relevant Land Acquisition Act sections, explains compensation calculation method, offers to draft an RTI application, and flags case to DLSA if complexity threshold is exceeded.

- **US-03: RTI Application Generation**

As a daily wage worker in Tamil Nadu whose MGNREGA wages have not been paid for 3 months, I want to generate a ready-to-file RTI application in Tamil, so that I can demand accountability from the local panchayat office.

Acceptance Criteria: DocumentAgent generates a correctly formatted RTI application with all mandatory fields, translates it to Tamil, and provides submission instructions.

- **US-04: Voice Input Accessibility**

As an elderly woman in rural Rajasthan with low literacy, I want to speak my legal question in Rajasthani Hindi into my phone, so that I can get help without needing to type or read.

Acceptance Criteria: IndicTrans2 transcribes voice input with >90% accuracy in Hindi dialects, agent processes and responds in the same language, response is also available as audio output.

- **US-05: Cross-Session Memory**

As a user who discussed a land dispute case with LegalSaathi last week, I want the agent to remember my case details when I return, so that I do not have to repeat my entire situation from scratch every time.

Acceptance Criteria: Mastra Memory persists case context across sessions, agent opens follow-up conversations with a summary of previous context, user can correct or update stored context.

6. Functional Requirements

ID	Requirement	Description
FR-01	Multilingual Voice Input	Accept voice queries in 10+ Indian languages via IndicTrans2 by AI4Bharat with >90% transcription accuracy
FR-02	Multilingual Text Input	Accept text queries in Hindi, Tamil, Telugu, Bengali, Marathi, Kannada, Gujarati, Malayalam, Punjabi, Odia
FR-03	Intent Classification	Mastra Intent Router Agent classifies every query into Scheme Eligibility, Legal Rights, or Document Generation within 500ms
FR-04	Semantic Legal Retrieval	Qdrant returns top-5 semantically relevant chunks from 4 legal collections for every agent query
FR-05	Scheme Eligibility Check	SchemeAgent checks eligibility for 500+ government schemes using MyScheme.gov.in API and Qdrant scheme collection
FR-06	Legal Rights Retrieval	LegalAgent retrieves relevant IPC, CrPC, BNS 2023 sections and state-specific land laws
FR-07	RTI Document Generation	DocumentAgent generates correctly formatted RTI applications and complaint letters in the user's language
FR-08	Cross-Session Memory	Mastra Memory persists user case context across conversations using Redis backing store
FR-09	Hallucination Detection	Enkrypt AI cross-checks every factual claim against retrieved source chunks before response delivery
FR-10	Confidence Scoring	Enkrypt AI assigns 0-1 confidence score; responses below 0.7 are blocked automatically
FR-11	DLSA Referral	Complex cases exceeding confidence threshold trigger automatic referral to nearest District Legal Services Authority
FR-12	PII Filtering	Enkrypt AI output guardrails detect and redact personally identifiable information before response delivery
FR-13	WhatsApp Integration	Full agent functionality available via WhatsApp Business API for users without smartphone app access
FR-14	Offline PWA Support	Progressive Web App functions with cached legal FAQs and basic guidance in zero-connectivity scenarios
FR-15	Document Scanner	Camera-based document capture sends government notices and legal documents for agent analysis

7. Non-Functional Requirements

Response Time	End-to-end query response in under 3 seconds (P95)
Uptime	99.5% availability — critical for users in legal emergencies
Languages Supported	10+ Indian languages via IndicTrans2
Hallucination Rate	Below 5% as measured by Enkrypt AI confidence scoring
Concurrent Users	Handles 10,000 concurrent users without degradation
Data Security	PII encrypted at rest (AES-256) and in transit (TLS 1.3)
Compliance	Adheres to IT Act 2000, DPDP Act 2023 for data handling
Accessibility	WCAG 2.1 AA compliant web interface
Mobile Performance	Core Web Vitals: LCP < 2.5s, CLS < 0.1 on 3G connections
Audit Trail	All Enkrypt AI evaluations logged with timestamps for compliance

8. Technical Architecture Summary

8.1 Technology Stack

Technology	Role	Specific Usage in LegalSaathi
Mastra	Agent Orchestration	Orchestrates 4 specialized agents, manages workflow pipeline, handles cross-session memory routing
Qdrant Cloud	Vector Memory & RAG	Stores 4 legal collections, handles semantic search, returns top-k chunks with source attribution
Enkrypt AI	Safety & Guardrails	Hallucination detection, confidence scoring, PII filtering, output evaluation pipeline
Next.js 14	Frontend	Web UI with i18n support for 10+ languages, PWA capabilities, document scanner integration
IndicTrans2 (AI4Bharat)	Multilingual NLP	Voice transcription and translation across 10+ Indian languages
Redis	Memory Store	Backing store for Mastra Memory — persists user case context across sessions
GPT-4o / Claude	LLM	Primary language model accessed via Mastra model provider abstraction layer
WhatsApp Business API	Messaging Channel	Enables agent access for users without smartphone app or web browser

9. Agent Architecture

9.1 Workflow Pipeline

Every user query follows a deterministic 7-step pipeline:

1. User sends query via Web, WhatsApp, or Voice in any supported Indian language
2. Next.js frontend / WhatsApp webhook passes query to Mastra Workflow Engine via REST API
3. Mastra Intent Router Agent classifies query into: Scheme Eligibility / Legal Rights / Document Generation
4. Relevant sub-agent (SchemeAgent / LegalAgent / DocumentAgent) is activated with case context from Mastra Memory
5. Sub-agent sends semantic search query to Qdrant and retrieves top-5 most relevant legal chunks
6. LLM generates a draft response grounded entirely in retrieved chunks
7. Enkrypt AI evaluates draft for hallucinations and confidence. Pass (≥ 0.7) → deliver to user. Fail (< 0.7) → block + explain + DLSA referral

9.2 Agent Responsibilities

Agent	Trigger Condition	Responsibilities
Intent Router Agent	Every incoming query	Language detection, intent classification, case context retrieval from Mastra Memory, sub-agent routing
SchemeAgent	Query classified as Scheme Eligibility	Check eligibility for 500+ schemes, query Qdrant scheme collection, call MyScheme.gov.in API, explain enrollment steps
LegalAgent	Query classified as Legal Rights	Retrieve IPC/CrPC/BNS 2023 sections, state land laws, fundamental rights, explain in plain language
DocumentAgent	Query classified as Document Generation	Generate RTI applications, complaint letters, legal notices using RTI Templates collection from Qdrant

9.3 Memory Architecture

Mastra Memory uses Redis as a backing store to persist the following per user:

- Case summary — compressed summary of the user's legal situation
- Query history — last 10 queries and responses for context
- Documents generated — RTI applications and complaints drafted in previous sessions
- Referral status — whether case has been escalated to DLSA

Memory is scoped per user ID, encrypted at rest using AES-256, and automatically expires after 90 days of inactivity.

9.4 Qdrant Collections

Collection	Content	Size
IPC & CrPC Collection	Indian Penal Code sections, CrPC procedures, BNS 2023 updates, criminal law precedents	~2,000 chunks
Government Schemes Collection	PM-KISAN, MGNREGA, PM Awas Yojana, 500+ scheme documents from MyScheme.gov.in	~5,000 chunks
Land Laws Collection	State-specific land acquisition acts, tenancy rights, forest rights, tribal land protection	~1,500 chunks
RTI Templates Collection	RTI application templates, successful precedents, complaint formats, DLSA contact directory	~800 chunks

Embedding model: text-embedding-3-small. Chunk size: 512 tokens with 20% overlap. Retrieval: Top-k = 5 per query.

10. Risks and Mitigations

Risk	Impact	Mitigation
Hallucinated legal advice	High — could lead to wrongful legal action by user	Enkrypt AI hallucination detection blocks responses below 0.7 confidence. All claims cross-checked against source chunks.
Language transcription errors	Medium — misunderstood query leads to wrong agent routing	IndicTrans2 fine-tuned on legal domain vocabulary. Fallback to text input if voice confidence is low.
Stale scheme data	Medium — outdated eligibility criteria misleads users	MyScheme.gov.in API integration ensures live scheme data. Qdrant collection updated monthly.
Complex case handled without lawyer	High — user acts on incomplete	Confidence-based DLSA redirect guardrail in Enkrypt AI. Disclaimer on every response recommending professional consultation for complex matters.

		advice in serious matter	
Data breach	privacy	High sensitive queries exposed — legal	PII filtering via Enkrypt AI. AES-256 encryption at rest. TLS 1.3 in transit. DPDP Act 2023 compliance.
Qdrant outdated	collection	Medium wrong section post BNS transition — IPC cited 2023	BNS 2023 already included in IPC collection. Quarterly update pipeline scheduled.

11. Project Timeline

Round 1 (June 23 – July 1, 2026)	Architecture design, agent workflow definition, Qdrant collection planning, Enkrypt AI integration design, PRD and Solution Overview Document submission
Round 2 — Sprint 1 (if qualified)	Mastra agent setup, Qdrant collection ingestion (IPC + Schemes), basic Intent Router working end-to-end
Round 2 — Sprint 2	LegalAgent + SchemeAgent fully functional with RAG, Enkrypt AI hallucination detection integrated
Round 2 — Sprint 3	DocumentAgent (RTI drafting), Mastra Memory (cross-session), WhatsApp integration
Round 2 — Sprint 4	Multilingual support (IndicTrans2), frontend polish, end-to-end testing, demo video recording
Demo Day	Live demo: farmer query in Hindi → agent routing → Qdrant retrieval → Enkrypt safety check → verified response